

Project Title:

Predictive Modeling for Breast Cancer Diagnosis Using Machine Learning

Feature Descriptions for Breast Cancer Dataset

- **id**: Unique identifier for each sample.
- **diagnosis**: Target variable indicating the diagnosis (**M** = Malignant, **B** = Benign).

Mean Features:

- **radius_mean**: Mean radius of the tumor cells.
- **texture_mean**: Mean texture (variation in gray levels) of the tumor cells.
- **perimeter_mean**: Mean perimeter of the tumor cells.
- **area_mean**: Mean area of the tumor cells.
- **smoothness_mean**: Mean smoothness (local variation in radius lengths) of the tumor cells.
- **compactness_mean**: Mean compactness ($\text{perimeter}^2 / \text{area} - 1.0$) of the tumor cells.
- **concavity_mean**: Mean concavity (severity of concave portions of the contour) of the tumor cells.
- **concave points_mean**: Mean number of concave portions of the tumor cell contours.
- **symmetry_mean**: Mean symmetry of the tumor cells.
- **fractal_dimension_mean**: Mean fractal dimension ("coastline approximation") of the tumor cells.

Standard Error Features:

- **radius_se**: Standard error of the radius of the tumor cells.
- **texture_se**: Standard error of the texture of the tumor cells.
- **perimeter_se**: Standard error of the perimeter of the tumor cells.
- **area_se**: Standard error of the area of the tumor cells.
- **smoothness_se**: Standard error of the smoothness of the tumor cells.
- **compactness_se**: Standard error of the compactness of the tumor cells.
- **concavity_se**: Standard error of the concavity of the tumor cells.
- **concave points_se**: Standard error of the number of concave portions of the tumor cell contours.
- **symmetry_se**: Standard error of the symmetry of the tumor cells.
- **fractal_dimension_se**: Standard error of the fractal dimension of the tumor cells.

Worst (Largest) Features:

- **radius_worst**: Largest (worst) radius of the tumor cells.
- **texture_worst**: Largest (worst) texture of the tumor cells.
- **perimeter_worst**: Largest (worst) perimeter of the tumor cells.
- **area_worst**: Largest (worst) area of the tumor cells.
- **smoothness_worst**: Largest (worst) smoothness of the tumor cells.
- **compactness_worst**: Largest (worst) compactness of the tumor cells.
- **concavity_worst**: Largest (worst) concavity of the tumor cells.
- **concave points_worst**: Largest (worst) number of concave portions of the tumor cell contours.
- **symmetry_worst**: Largest (worst) symmetry of the tumor cells.
- **fractal_dimension_worst**: Largest (worst) fractal dimension of the tumor cells.

Import Libraries

```
In [4]: # Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, f1_score, classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")
```

Load the Cancer Wisconsin dataset

```
In [6]: df = pd.read_csv('Cancer Wisconsin.csv')
```

```
In [7]: df.head()
```

```
Out[7]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavit
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	

5 rows × 33 columns



Information of dataset

```
In [9]: df.shape
```

```
Out[9]: (569, 33)
```

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                 569 non-null    float64
12  radius_se                              569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                                569 non-null    float64
16  smoothness_se                          569 non-null    float64
17  compactness_se                         569 non-null    float64
18  concavity_se                           569 non-null    float64
19  concave points_se                      569 non-null    float64
20  symmetry_se                            569 non-null    float64
21  fractal_dimension_se                   569 non-null    float64
22  radius_worst                           569 non-null    float64
23  texture_worst                          569 non-null    float64
24  perimeter_worst                        569 non-null    float64
25  area_worst                             569 non-null    float64
26  smoothness_worst                       569 non-null    float64
27  compactness_worst                      569 non-null    float64
28  concavity_worst                        569 non-null    float64
29  concave points_worst                   569 non-null    float64
30  symmetry_worst                         569 non-null    float64
31  fractal_dimension_worst                 569 non-null    float64
32  Unnamed: 32                             0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

Check The Column Names

```
In [12]: df.columns
```

```
Out[12]: Index(['id', 'diagnosis', 'radius_mean', 'texture_mean', 'perimeter_mean',
              'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
              'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
              'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
              'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
              'fractal_dimension_se', 'radius_worst', 'texture_worst',
              'perimeter_worst', 'area_worst', 'smoothness_worst',
              'compactness_worst', 'concavity_worst', 'concave points_worst',
              'symmetry_worst', 'fractal_dimension_worst', 'Unnamed: 32'],
              dtype='object')
```

Rename The Columns

```
In [20]: df = df.rename(columns={
          'concave points_mean': 'concave_points_mean',
          'concave points_worst': 'concave_points_worst',
        })
```

Data Cleaning

```
In [23]: # Check for missing values
print("\nMissing Values:\n", df.isnull().sum())
```

```
Missing Values:
id                0
diagnosis         0
radius_mean       0
texture_mean      0
perimeter_mean    0
area_mean         0
smoothness_mean   0
compactness_mean  0
concavity_mean    0
concave_points_mean 0
symmetry_mean     0
fractal_dimension_mean 0
radius_se         0
texture_se        0
perimeter_se      0
area_se           0
smoothness_se     0
compactness_se    0
concavity_se      0
concave points_se 0
symmetry_se       0
fractal_dimension_se 0
radius_worst      0
texture_worst     0
perimeter_worst   0
area_worst        0
smoothness_worst  0
compactness_worst 0
concavity_worst   0
concave_points_worst 0
symmetry_worst    0
fractal_dimension_worst 0
Unnamed: 32       569
dtype: int64
```

```
In [25]: # Drop Unwanted Columns
df.drop(['id', 'Unnamed: 32'], axis=1, inplace=True)
```

```
In [27]: # Check the balance of the target classes
df['diagnosis'].value_counts()
```

```
Out[27]: diagnosis
B      357
M      212
Name: count, dtype: int64
```

```
In [29]: # Change The Diagnosis in Numeric (M=1, B=0)
df['diagnosis'] = df['diagnosis'].map({'M':1, 'B':0})
```

Check The Duplicate Values

```
In [32]: # check the duplicate values
df.duplicated().sum()
```

Out[32]: 0

Summary Statistics

In [35]: df.describe().T

Out[35]:


	count	mean	std	min	25%	50%	75%	max
diagnosis	569.0	0.372583	0.483918	0.000000	0.000000	0.000000	1.000000	1.00000
radius_mean	569.0	14.127292	3.524049	6.981000	11.700000	13.370000	15.780000	28.11000
texture_mean	569.0	19.289649	4.301036	9.710000	16.170000	18.840000	21.800000	39.28000
perimeter_mean	569.0	91.969033	24.298981	43.790000	75.170000	86.240000	104.100000	188.50000
area_mean	569.0	654.889104	351.914129	143.500000	420.300000	551.100000	782.700000	2501.00000
smoothness_mean	569.0	0.096360	0.014064	0.052630	0.086370	0.095870	0.105300	0.16340
compactness_mean	569.0	0.104341	0.052813	0.019380	0.064920	0.092630	0.130400	0.34540
concavity_mean	569.0	0.088799	0.079720	0.000000	0.029560	0.061540	0.130700	0.42680
concave_points_mean	569.0	0.048919	0.038803	0.000000	0.020310	0.033500	0.074000	0.20120
symmetry_mean	569.0	0.181162	0.027414	0.106000	0.161900	0.179200	0.195700	0.30400
fractal_dimension_mean	569.0	0.062798	0.007060	0.049960	0.057700	0.061540	0.066120	0.09744
radius_se	569.0	0.405172	0.277313	0.111500	0.232400	0.324200	0.478900	2.87300
texture_se	569.0	1.216853	0.551648	0.360200	0.833900	1.108000	1.474000	4.88500
perimeter_se	569.0	2.866059	2.021855	0.757000	1.606000	2.287000	3.357000	21.98000
area_se	569.0	40.337079	45.491006	6.802000	17.850000	24.530000	45.190000	542.20000
smoothness_se	569.0	0.007041	0.003003	0.001713	0.005169	0.006380	0.008146	0.03113
compactness_se	569.0	0.025478	0.017908	0.002252	0.013080	0.020450	0.032450	0.13540
concavity_se	569.0	0.031894	0.030186	0.000000	0.015090	0.025890	0.042050	0.39600
concave_points_se	569.0	0.011796	0.006170	0.000000	0.007638	0.010930	0.014710	0.05279
symmetry_se	569.0	0.020542	0.008266	0.007882	0.015160	0.018730	0.023480	0.07895
fractal_dimension_se	569.0	0.003795	0.002646	0.000895	0.002248	0.003187	0.004558	0.02984
radius_worst	569.0	16.269190	4.833242	7.930000	13.010000	14.970000	18.790000	36.04000
texture_worst	569.0	25.677223	6.146258	12.020000	21.080000	25.410000	29.720000	49.54000
perimeter_worst	569.0	107.261213	33.602542	50.410000	84.110000	97.660000	125.400000	251.20000
area_worst	569.0	880.583128	569.356993	185.200000	515.300000	686.500000	1084.000000	4254.00000
smoothness_worst	569.0	0.132369	0.022832	0.071170	0.116600	0.131300	0.146000	0.22260
compactness_worst	569.0	0.254265	0.157336	0.027290	0.147200	0.211900	0.339100	1.05800
concavity_worst	569.0	0.272188	0.208624	0.000000	0.114500	0.226700	0.382900	1.25200
concave_points_worst	569.0	0.114606	0.065732	0.000000	0.064930	0.099930	0.161400	0.29100
symmetry_worst	569.0	0.290076	0.061867	0.156500	0.250400	0.282200	0.317900	0.66380
fractal_dimension_worst	569.0	0.083946	0.018061	0.055040	0.071460	0.080040	0.092080	0.20750

In []:

Data Exploration


In [39]:

```
# 4.1 Diagnosis Count Plot
plt.figure(figsize=(10,8))
sns.countplot(x='diagnosis', data=df)
plt.title("Benign (0) vs Malignant (1) Cases")
plt.show()
```



Correlation Matrix

```
In [42]: plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), annot=False, cmap='coolwarm')
plt.title("Features Correlation")
plt.show()
```



Define independent variables (X) and dependent variable (Y)

```
In [45]: # Features (X) and Target (y)
X = df.drop('diagnosis', axis=1)
y = df['diagnosis']
```

Feature Scaling

```
In [48]: # Feature Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Split Data into Training & Testing Sets

```
In [51]: # Train-Test Split (80% Train, 20% Test)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

Defining & Training All The Model For Choosing the Best Model

```
In [54]: # Test Of Multiple Models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "Support Vector Machine": SVC(),
    "Naive Bayes": GaussianNB(),
```

```
}
```

```
In [56]: results = []
for name, model in models.items():
    # Model Training
    model.fit(X_train, y_train)

    # For Predictions
    y_pred = model.predict(X_test)

    # Check Performance
    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    print(f"*60)
    results.append([name, acc, f1])

    # Classification Report
    print(f"\nModel: {name} \n")
    print("Accuracy:", round(acc, 4))
    print("F1 Score:", round(f1, 4))
    print(f"*60)
    print("Classification Report:\n", classification_report(y_test, y_pred))
```

Model: Logistic Regression

Accuracy: 0.9737

F1 Score: 0.9647

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	71
1	0.98	0.95	0.96	43
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

Model: Decision Tree

Accuracy: 0.9386

F1 Score: 0.9176

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.96	0.95	71
1	0.93	0.91	0.92	43
accuracy			0.94	114
macro avg	0.94	0.93	0.93	114
weighted avg	0.94	0.94	0.94	114

Model: Random Forest

Accuracy: 0.9649

F1 Score: 0.9524

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.99	0.97	71
1	0.98	0.93	0.95	43
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

Model: Support Vector Machine

Accuracy: 0.9737

F1 Score: 0.9647

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.97	0.99	0.98	71
1	0.98	0.95	0.96	43
accuracy			0.97	114
macro avg	0.97	0.97	0.97	114
weighted avg	0.97	0.97	0.97	114

Model: Naive Bayes

Accuracy: 0.9649

F1 Score: 0.9524

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.99	0.97	71
1	0.98	0.93	0.95	43
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

In []:

Drop Highly Correlated Features

Calculate correlation matrix

```
In [61]: # Calculate correlation matrix
corr_matrix = X.corr().abs()
```

Select upper triangle of correlation matrix

```
In [64]: # Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
```

Find features with correlation > 0.95

```
In [67]: to_drop = [column for column in upper.columns if any(upper[column] > 0.95)]
print("Highly correlated features:", to_drop)
```

Highly correlated features: ['perimeter_mean', 'area_mean', 'perimeter_se', 'area_se', 'radius_worst', 'perimeter_worst', 'area_worst']

Remove Multicollinearity

```
In [70]: # Drop them
X_reduced = X.drop(to_drop, axis=1)
```

Display Summary Of Results

```
In [73]: import statsmodels.api as sm

X_with_const = sm.add_constant(X_reduced)
logit_model = sm.Logit(y, X_with_const).fit()
print(logit_model.summary())
```


Optimization terminated successfully.
Current function value: 0.036903
Iterations 23

Logit Regression Results

Dep. Variable:	diagnosis	No. Observations:	569
Model:	Logit	Df Residuals:	545
Method:	MLE	Df Model:	23
Date:	Wed, 16 Apr 2025	Pseudo R-squ.:	0.9441
Time:	11:49:56	Log-Likelihood:	-20.998
converged:	True	LL-Null:	-375.72
Covariance Type:	nonrobust	LLR p-value:	4.545e-135

	coef	std err	z	P> z	[0.025	0.975]
const	-103.8447	36.003	-2.884	0.004	-174.409	-33.280
radius_mean	1.7697	0.791	2.238	0.025	0.220	3.320
texture_mean	-0.2662	0.334	-0.797	0.425	-0.921	0.388
smoothness_mean	182.4330	118.254	1.543	0.123	-49.341	414.207
compactness_mean	-138.6832	93.874	-1.477	0.140	-322.673	45.306
concavity_mean	107.8974	68.669	1.571	0.116	-26.692	242.487
concave_points_mean	3.5654	108.139	0.033	0.974	-208.384	215.514
symmetry_mean	-55.3856	42.870	-1.292	0.196	-139.409	28.638
fractal_dimension_mean	124.2543	315.092	0.394	0.693	-493.315	741.824
radius_se	38.8085	13.128	2.956	0.003	13.077	64.540
texture_se	-5.5030	2.349	-2.342	0.019	-10.108	-0.898
smoothness_se	605.4090	350.276	1.728	0.084	-81.120	1291.938
compactness_se	333.3153	177.619	1.877	0.061	-14.812	681.443
concavity_se	-202.0945	93.101	-2.171	0.030	-384.568	-19.621
concave_points_se	676.0393	420.760	1.607	0.108	-148.636	1500.714
symmetry_se	-264.4777	187.604	-1.410	0.159	-632.174	103.218
fractal_dimension_se	-4185.1475	1601.426	-2.613	0.009	-7323.884	-1046.411
texture_worst	0.9795	0.352	2.782	0.005	0.289	1.669
smoothness_worst	-53.4620	63.474	-0.842	0.400	-177.868	70.944
compactness_worst	-43.6258	31.653	-1.378	0.168	-105.665	18.414
concavity_worst	22.3973	18.855	1.188	0.235	-14.559	59.353
concave_points_worst	33.7196	60.301	0.559	0.576	-84.468	151.907
symmetry_worst	53.8119	27.327	1.969	0.049	0.252	107.372
fractal_dimension_worst	461.8041	206.469	2.237	0.025	57.132	866.477

Possibly complete quasi-separation: A fraction 0.77 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

- The model is indeed showing good accuracy; however, analysis revealed that out of the 31 features, only 8 to 10 are truly significant. This indicates that the model is also utilizing some irrelevant or noisy features, which could lead to overfitting. In the next phase, I plan to apply feature selection techniques to optimize the model and improve its interpretability.

Feature Selection: Using Only Significant Features for Overfitting-Free Model

```
In [86]: # Features (X) and Target (y)
X = df[['radius_mean',
        'radius_se',
        'texture_se',
        'concavity_se',
        'fractal_dimension_se',
        'texture_worst',
        'symmetry_worst',
        'fractal_dimension_worst']]

y = df['diagnosis']

# Feature Scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-Test Split (80% Train, 20% Test)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

Again Check And Drop Highly Correlated Features

```
In [90]: # Calculate correlation matrix
corr_matrix = X.corr().abs()

# Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))

# Find features with correlation > 0.95
to_drop = [column for column in upper.columns if any(upper[column] > 0.95)]

print("Highly correlated features:", to_drop)

# Drop them
X_reduced = X.drop(to_drop, axis=1)
```

Highly correlated features: []

Again Check And Display Summary Of Results

```
In [93]: import statsmodels.api as sm

X_with_const = sm.add_constant(X_reduced)
logit_model = sm.Logit(y, X_with_const).fit()
print(logit_model.summary())
```

Optimization terminated successfully.

Current function value: 0.083104

Iterations 12

```
Logit Regression Results
=====
Dep. Variable:      diagnosis    No. Observations:      569
Model:              Logit       Df Residuals:          560
Method:             MLE        Df Model:              8
Date:               Wed, 16 Apr 2025    Pseudo R-squ.:      0.8741
Time:               11:50:29    Log-Likelihood:     -47.286
converged:          True        LL-Null:            -375.72
Covariance Type:    nonrobust    LLR p-value:        1.375e-136
=====
                    coef    std err          z      P>|z|      [0.025      0.975]
-----
const              -47.7101      6.429     -7.421     0.000     -60.311     -35.109
radius_mean         1.3359      0.224      5.971     0.000      0.897      1.774
radius_se          18.1385      4.142      4.379     0.000     10.020     26.257
texture_se         -1.9022      1.005     -1.892     0.058     -3.872      0.068
concavity_se       32.0266     10.360      3.092     0.002     11.722     52.331
fractal_dimension_se -889.7916    280.682     -3.170     0.002    -1439.917    -339.666
texture_worst        0.4185      0.088      4.778     0.000      0.247      0.590
symmetry_worst      12.9287      5.841      2.214     0.027      1.481     24.376
fractal_dimension_worst 134.2854     31.808      4.222     0.000     71.944     196.627
=====
```

Possibly complete quasi-separation: A fraction 0.36 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

In []:

Results Comparison

```
In [96]: # Summary table
# Results Comparison
results_df = pd.DataFrame(results, columns=['Model', 'Accuracy', 'F1 Score'])
print("\nModels Comparison:")
results_df.sort_values('F1 Score', ascending=False)
```

Models Comparison:

```
Out[96]:
```

	Model	Accuracy	F1 Score
0	Logistic Regression	0.973684	0.964706
3	Support Vector Machine	0.973684	0.964706
2	Random Forest	0.964912	0.952381
4	Naive Bayes	0.964912	0.952381
1	Decision Tree	0.938596	0.917647


Model Performance

```
In [99]: import seaborn as sns
# Plot Model Performance
plt.figure(figsize=(10, 5))
ax = sns.barplot(x="F1 Score", y="Model", data=results_df, palette="Blues")

# Add annotations (F1 Score on bars)
for container in ax.containers:
    ax.bar_label(container, fmt="%.3f", fontsize=10, color="black", padding=1)

# Set title and limits
plt.title("Model Performance Comparison", fontsize=14, fontweight="bold")
plt.xlim(0, 1) #F1 Score range

# Show plot
plt.show()
```

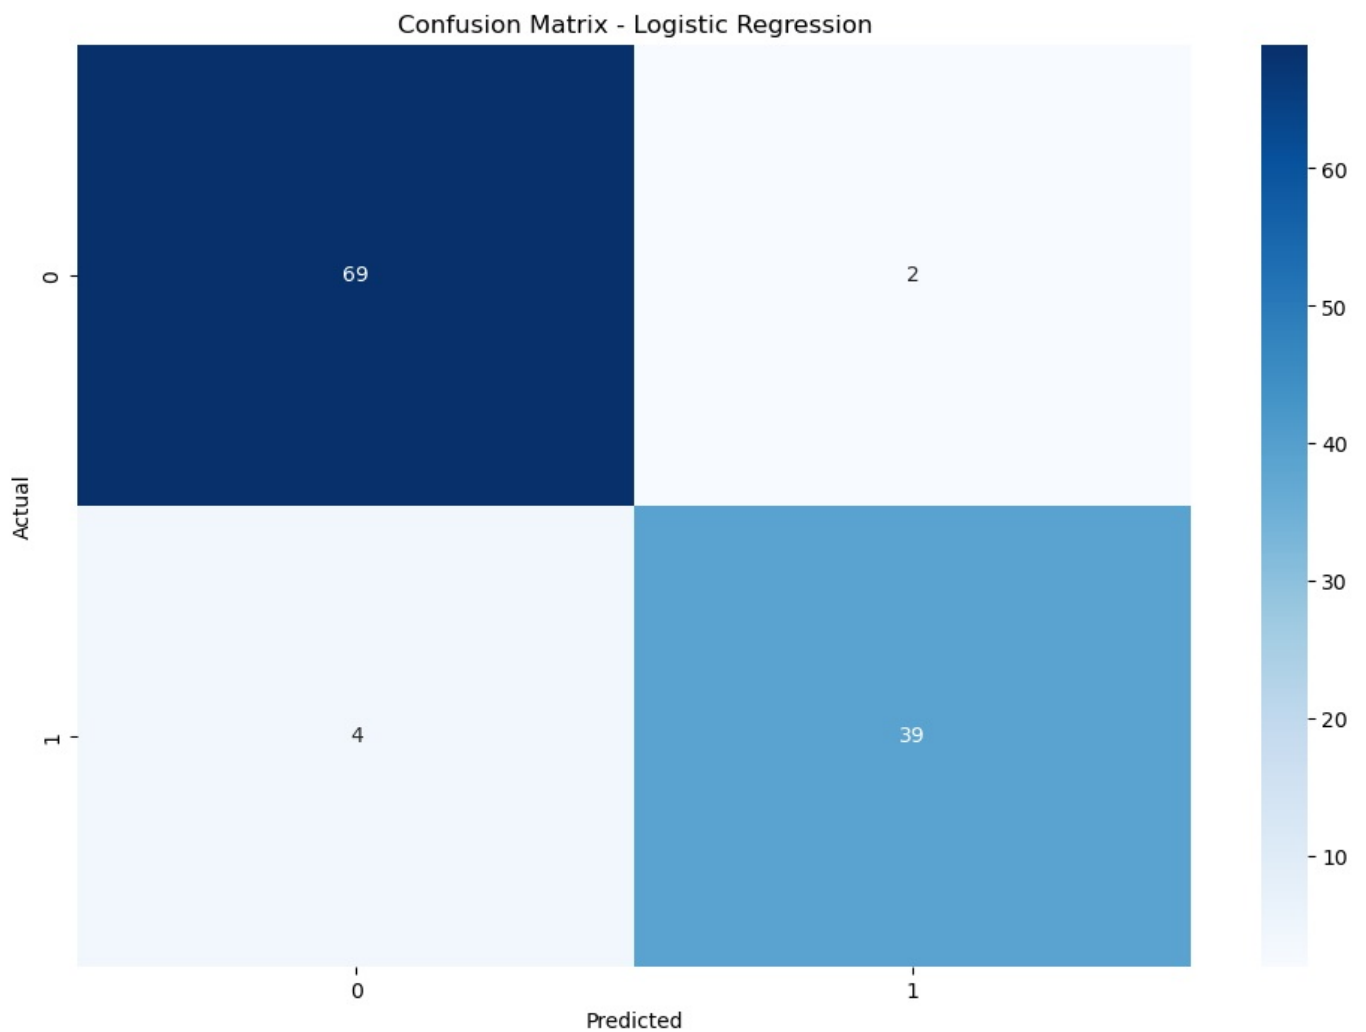


In []:

Best Model Of Confusion Matrix

```
In [102]: # Best Model Of Confusion Matrix
plt.figure(figsize=(12, 8))
best_model = LogisticRegression(max_iter=1000)
best_model.fit(X_train, y_train)
y_pred = best_model.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Logistic Regression")
plt.show()
```



Choose The Best Model

- Support Vector Machine And Logistic Regression Model is the best model for batter prediction

In []:

Prediction

Actual Data

In [105...

```
df[['radius_mean',
    'radius_se',
    'texture_se',
    'concavity_se',
    'fractal_dimension_se',
    'texture_worst',
    'symmetry_worst',
    'fractal_dimension_worst',
    'diagnosis']].tail(2)
```

Out[105]:

	radius_mean	radius_se	texture_se	concavity_se	fractal_dimension_se	texture_worst	symmetry_worst	fractal_dimension_wo
567	20.60	0.7260	1.595	0.07117	0.006185	39.42	0.4087	0.124
568	7.76	0.3857	1.428	0.00000	0.002783	30.37	0.2871	0.070

In []:

No Cancer data

In [107...

[illegible]

```
'fractal_dimension_se': 0.002783,  
'texture_worst': 30.37,  
'symmetry_worst': 0.2871,  
'fractal_dimension_worst': 0.07039  
]])
```

In [109.. new_patient_no_cancer

Out[109..

	radius_mean	radius_se	texture_se	concavity_se	fractal_dimension_se	texture_worst	symmetry_worst	fractal_dimension_wors
0	7.76	0.3857	1.428	0.0	0.002783	30.37	0.2871	0.07039

In [111..

```
# Feature Scaling for New Data  
new_patient_scaled = scaler.transform(new_patient_no_cancer)  
  
# Prediction  
prediction = best_model.predict(new_patient_scaled)  
prediction_proba = best_model.predict_proba(new_patient_scaled)  
  
# Result Display  
print("\nNew Patient Prediction:")  
print("Predicted Class:", "Malignant (Cancer)" if prediction[0] == 1 else "Benign (No Cancer)")  
print("Probability [Benign, Malignant]:", prediction_proba[0])
```

New Patient Prediction:
Predicted Class: Benign (No Cancer)
Probability [Benign, Malignant]: [0.99838012 0.00161988]

In []:

In []:

Cancer data

In [114..

```
new_patient= pd.DataFrame ([[ 'radius_mean': 20.60,  
                               'radius_se': 0.7260,  
                               'texture_se': 1.5950,  
                               'concavity_se': 0.07117,  
                               'fractal_dimension_se': 0.006185,  
                               'texture_worst': 39.42,  
                               'symmetry_worst': 0.4087,  
                               'fractal_dimension_worst': 0.12400  
                               ]])
```

In [116.. new_patient

Out[116..

	radius_mean	radius_se	texture_se	concavity_se	fractal_dimension_se	texture_worst	symmetry_worst	fractal_dimension_wors
0	20.6	0.726	1.595	0.07117	0.006185	39.42	0.4087	0.12400

In [118..

```
# Feature Scaling for New Data  
new_patient_scaled = scaler.transform(new_patient)  
  
# Prediction  
prediction = best_model.predict(new_patient_scaled)  
prediction_proba = best_model.predict_proba(new_patient_scaled)  
  
# Result Display  
print("\nNew Patient Prediction:")  
print("Predicted Class:", "Malignant (Cancer)" if prediction[0] == 1 else "Benign (No Cancer)")  
print("Probability [Benign, Malignant]:", prediction_proba[0])
```

New Patient Prediction:
Predicted Class: Malignant (Cancer)
Probability [Benign, Malignant]: [1.93916098e-07 9.99999806e-01]

In []:

In []:

Run The Model in Streamlit Web App

In []:

In [122..

```
import pandas as pd  
import numpy as np  
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.linear_model import LogisticRegression
import pickle

# Select the 8 features used for prediction
X = df[['radius_mean', 'radius_se', 'texture_se', 'concavity_se', 'fractal_dimension_se',
        'texture_worst', 'symmetry_worst', 'fractal_dimension_worst']]
y = df['diagnosis'] # Assuming 'diagnosis' is the target column

# Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train the model (Logistic Regression as an example)
model = LogisticRegression()
model.fit(X_scaled, y)

# Save the trained model and scaler
with open('cancer_model.pkl', 'wb') as model_file:
    pickle.dump(model, model_file)

with open('scaler.pkl', 'wb') as scaler_file:
    pickle.dump(scaler, scaler_file)

print("Model and scaler saved!")

```

Model and scaler saved!

In []:

```

In [125.%%writefile app.py
import streamlit as st
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
import pickle

# Load saved model and scaler
with open('cancer_model.pkl', 'rb') as model_file:
    model = pickle.load(model_file)
with open('scaler.pkl', 'rb') as scaler_file:
    scaler = pickle.load(scaler_file)

# App title and description
st.title("Breast Cancer Prediction App")
st.write("""
This app predicts whether a breast tumor is Malignant (Cancerous) or Benign (Non-Cancerous)
using machine learning. Enter the patient's details below:
""")

# Input form
st.header("Patient Details")
with st.form("prediction_form"):
    # Create input fields for the 8 features
    col1, col2 = st.columns(2)

    with col1:
        radius_mean = st.number_input("Radius Mean", min_value=0.0, value=7.76)
        radius_se = st.number_input("Radius SE", min_value=0.0, value=0.3857)
        texture_se = st.number_input("Texture SE", min_value=0.0, value=1.4280)
        concavity_se = st.number_input("Concavity SE", min_value=0.0, value=0.00000)

    with col2:
        texture_worst = st.number_input("Texture Worst", min_value=0.0, value=30.37)
        symmetry_worst = st.number_input("Symmetry Worst", min_value=0.0, value=0.2871)
        fractal_dimension_worst = st.number_input("Fractal Dimension Worst", min_value=0.0, value=0.07039)
        fractal_dimension_se = st.number_input("Fractal Dimension SE", min_value=0.0, value=0.002783)

    submit_button = st.form_submit_button("Predict Diagnosis")

# Prediction logic
if submit_button:
    # Create feature array with only the 8 selected features
    features = np.array([
        radius_mean, radius_se, texture_se, concavity_se, fractal_dimension_se,
        texture_worst, symmetry_worst, fractal_dimension_worst
    ])

    # Scale features
    features_scaled = scaler.transform(features)

    # Make prediction
    prediction = model.predict(features_scaled)

```

```

probability = model.predict_proba(features_scaled)

# Display results
st.header("Prediction Results")
if prediction[0] == 1:
    st.error(f" **Prediction:** Malignant (Cancerous) - {probability[0][1]*100:.2f}% probability")
else:
    st.success(f"✔ **Prediction:** Benign (Non-Cancerous) - {probability[0][0]*100:.2f}% probability")

# Show probability breakdown
st.write(f"**Probability Breakdown:**")
st.write(f"- Benign: {probability[0][0]*100:.2f}%")
st.write(f"- Malignant: {probability[0][1]*100:.2f}%")

# Run instructions
st.sidebar.header("How to Use")
st.sidebar.write("""
1. Enter patient's tumor characteristics
2. Click 'Predict Diagnosis'
3. View results
""")

```

Overwriting app.py

```

In [127... import subprocess
import sys

# Install streamlit if not installed
subprocess.check_call([sys.executable, "-m", "pip", "install", "streamlit"])

# Run the streamlit app
subprocess.Popen([sys.executable, "-m", "streamlit", "run", "app.py"])

```

Out[127... <Popen: returncode: None args: ['C:\\ProgramData\\anaconda3\\python.exe', '-...>